

## Data Engineering Take-Home Exercise

[PriceHubble](#) helps B2B clients make smart real estate decisions, such as evaluating property values, tracking market trends, spotting investment opportunities, and improving property management.

As a Senior Data Engineer in the Property Intelligence Tribe, you'll have a big impact by gathering and combining property listing and transaction data from different sources. You'll ensure data quality and enhance it with details like location and market trends to give a complete view of the residential property market. This data helps our clients with property valuation decisions and building trusted Automated Valuation Models (AVMs) for assessing property values.

This take-home project is designed to give you a feel for the kind of work you'd be doing at PriceHubble. We hope you enjoy it!

### Objective:

Your task is to build a data pipeline using the provided [JSONL](#) file. You'll process the data and insert it into a [DuckDB](#) table. You can choose any tech stack, as long as the code is reproducible and can be run by the reviewers.

### Requirements:

#### Programming Language:

- The code should be written in Python.

#### Input:

- You will receive a JSONL file with raw data (***scraping\_data.jsonl***).
- The input file contains one property offer per row, in JSON format. Each row has the following columns:

Column	Type	Description
id	string	Unique ID of the property
raw_price	string	Price info (e.g., "530 000€/mo.")
living_area	float	The area of the property in square meters

property_type	string	Type of property (e.g., house, studio)
municipality	string	City or town where the property is located
scraping_date	string	Date the data was scraped (YYYY-MM-DD)

### Example input:

Python

```
{
  "id": "0000a4fb",
  "raw_price": "530 000€/mo.",
  "living_area": 84.0,
  "property_type": "apartment",
  "municipality": "Solothurn",
  "scraping_date": "2021-02-17"
}
```

### DuckDB Table:

- Create a DuckDB table matching the output data structure.
- Ensure the table is populated correctly after running the pipeline.

The output should match the following structure:

Column	Type	Null able	Description
id	string	No	Property ID
scraping_date	string	No	Date when the data was scraped
property_type	string	No	Type of property
municipality	string	No	Municipality of the property
price	float	No	Converted price in numeric format

living_area	float	No	Area of the property
price_per_square_meter	float	No	Price per square meter

### Example output:

Python

```
{
  "id": "0000a4fb",
  "scraping_date": "2021-02-17",
  "property_type": "apartment",
  "municipality": "Solothurn",
  "price": 530000.0,
  "living_area": 84.0,
  "price_per_square_meter": 6309.52
}
```

### Filtering Criteria:

- Only include rows where:
  - `price_per_square_meter` is between 500 and 15,000.
  - `property_type` is either "apartment" or "house".
  - `scraping_date` is after March 5, 2020.

### Orchestration:

- Use any orchestration tool of your choice (e.g., Airflow, Prefect, Dagster, Meltano).

### Reproducibility:

- Ensure the pipeline can be easily run by the reviewers.
- Use Docker (or an alternative solution) to set up a reproducible environment.
- Provide clear instructions on setting up and running the pipeline, including dependencies.

### Deliverables:

#### Code:

- Complete Python code for the data pipeline, including extraction, loading, and transformation steps.
- Any necessary setup scripts (e.g., Dockerfiles, configuration files).
- Orchestration files (e.g., DAGs or Prefect flows).

**Documentation:**

- A README with clear instructions on setting up and running the pipeline:
  - Prerequisites (e.g., Docker, libraries, tools).
  - How to run the pipeline manually or through an orchestrator.
  - Example outputs or logs from the pipeline.

**Evaluation Criteria:**

1. **Code Quality:**
  - Clear, well-structured Python code following best practices.
2. **Reproducibility:**
  - The pipeline should be easy to set up and run without complex configurations.
3. **Scalability:**
  - The solution should handle large datasets effectively.
4. **Documentation:**
  - Instructions should be clear and easy to follow.

**Additional Info:**

- You're free to use any libraries or tools you're comfortable with.
- Docker is recommended to ensure the environment is reproducible.

**Time Limit:**

- Expect to spend 4 to 6 hours. Focus on simplicity and clarity.